



University
Mohammed VI
Polytechnic



Deliverable #: Title of Lab/Project here (e.g.: Conceptual Design)

Data Management Course
UM6P College of Computing

Professor: Karima Echihabi **Program:** Computer Engineering
Session: Fall 2025

Team Information

Team Name
Member 1	Adam Chikri
Member 2	Amine Abdellaoui
Member 3	Douae Bouayadi
Member 4	Aya Abid
Member 5	Nissrine Benallal
Member 6	Yasmin Bouaboud
Repository Link	https://github.com/Adam-chi/Lab5-DBMS

1 Introduction

This lab focused on designing and implementing a hospital management database. Tasks included creating and normalizing tables, enforcing data integrity with keys and constraints, inserting and manipulating sample data, and writing SQL queries to extract meaningful information. The work provided hands-on experience in building a structured, reliable, and database system that we can run queries on.

2 Requirements

The objective of this project part was to work through several tasks related to refining and operating the MNHS database .We first hade to validate each relation against BCNF and verify lossless joins and dependency preservation after decomposition.After that ,we implemented the corrected shema using SQL DDL,populated it with sample data using DML operations .Finally , we executed the 20 queries on the MNHS schema created .

3 Methodology

This part will provide the logic when working with BCNF: Validating the BCNF, and Testing the lossless joins, and finally checking if it's dependency preserving.

A. BCNF Validation

For each relation (Patient, ContactLocation, Staff, Hospital, Department, Insurance, ClinicalActivity, Expense, Appointment, Emergency, Medication, Stock, Prescription, Includes, have, Patient–Insurance, Work_in). The first step is listing the FDs then Identifying keys and by the end verifying superkey condition.

All relations were found to be already in BCNF because all FDs had a superkey on the left side.

B. Lossless Join Verification

To test lossless join, we used the standard condition:

For a decomposition $R \rightarrow R_1$ and R_2 , the join is lossless if:

$$R_1 \cap R_2 \rightarrow R_1 \quad \text{or} \quad R_1 \cap R_2 \rightarrow R_2$$

In other words, the common attributes between the decomposed relations must form a key for at least one of them.

For each relation in our schema that required decomposition, we documented the intersection and verified the lossless condition.

C. Dependency Preservation Verification

After decomposition, we checked whether all functional dependencies of the original relation were still preserved.

Compute the union of FDs in decomposed relations

First step: Compute the closure:

$$(F'_1 \cup F'_2 \cup \dots \cup F'_n)^+$$

Second step: Compare with the closure of the original FDs.

If we have:

$$(F'_1 \cup F'_2 \cup \dots \cup F'_n)^+ = F^+$$

then the dependency preservation is satisfied.

Using closure-based reasoning, we confirmed: Patient: all FDs preserved in R_1 and R_2

Conclusion: All decompositions preserve functional dependencies.

This methodology part will be dedicated to explain our logic behind the queries.

Query 1:

we order patients by Fullname (there is no attribute called last name) using ORDER BY FullName

Query 2:

This query lists all distinct insurance types available in the system. By using SELECT DISTINCT on the Insurance table, it removes duplicates and provides a clear overview of coverage options offered to patients, useful for administrative reference and patient registration.

Query 3:

This query retrieves the ID's and full names of staff members who work in hospitals located in Rabat. To do that, we join the Staff, work_in, Department, and Hospital tables to connect staff with hospitals through their department assignments, we then filter the results to only include hospitals where the city is Rabat. DISTINCT is used to avoid duplicate entries if staff work in multiple departments or hospitals in Rabat.

Query 4:

To get the appointments scheduled for the next 7 days, we'd need to join Appointments with Clinical Activity since it's the table that has the dates and then retrieving the rows where the date is between today and after 7 days.

Query 5:

The Appointment and ClinicalActivity tables were joined to link each appointment to its corresponding department. The results were then grouped by department, and the number of appointments in each group was calculated to analyze their distribution

Query 6:

We calculate the average unit price using AVG(UnitPrice) from the table Stock

Query 7:

This query identifies hospitals with high emergency admission rates. By joining the Hospital, Department, ClinicalActivity, and Emergency tables, it connects each hospital to its departments and emergency clinical activities. The filter on 'Admitted' ensures only serious cases are counted, and the HAVING clause isolates facilities with more than 20 admissions, highlighting those that might need additional resources.

Query 8:

To get the unit price of a medication we'll need to join the medication table with the stock, and then comparing the rows to see where the unit price is below 200 and in the 'Antibiotic' Therapeutic Class

Query 9:

Keep medication if there are <3 medications in the same hospital that cost more.

Query 10:

This query returns the number of scheduled, completed, and cancelled appointments for each department, all in one result table. Here's how it works: DEP_ID and department name are selected from Department, then three separate subqueries are used to count scheduled, cancelled and completed appointments. Each one of the three joins ClinicalActivity and Appointment, filtering by department and status. After that counts are arranged as new columns for each department in the output.

Query 11:

To identify patients with no scheduled appointments in the next 30 days, a subquery was used to find patients with upcoming scheduled appointments, and then only those patients whose IDs were not in this list were selected. This highlighted patients who are not scheduled for any consultations in the near future

Query 12:

This query calculates for each staff member the total number of appointments they are responsible for, and their percentage share of all appointments in their hospital. It joins Staff, ClinicalActivity, Appointment, and Department to relate each staff member to their department and hospital. It then counts appointments per staff (COUNT(A.CAID)), the percentage is calculated as (staff member's appointment count) divided by (total appointments in their hospital), times 100. It uses a subquery in the percentage calculation to get hospital totals. The results are grouped by staff to display statistics for each individual.

Query 13:

To monitor hospital stock and identify medications below the reorder level, the Stock table was queried to find entries where the current quantity was less than the reorder threshold. Joining with the Hospital and Medication tables allowed for more descriptive output, showing which medications in which hospitals needed attention

Query 14:

We used the basic division pattern seen in the lecture by changing the table names to get the appropriate result, a logical way to see this, is if for a hospital the table of the total antibiotic medications that exist minus the antibiotics that the hospital has is empty, then the hospital has every single antibiotic in the catalog.

Query 15:

This query analyzes medication pricing strategies between hospitals in the same city. It first calculates the average price per hospital and therapeutic class through joins between Hospital, Stock, and Medication. Correlated subqueries then compute average prices by city and drug class. Finally, UNION combines two sets: hospitals with prices above the local average and those with prices below or equal to average, providing a complete picture of pricing positioning.

Query 16:

This query determines the next appointment date for each patient. By linking the Patient, ClinicalActivity, and Appointment tables, it identifies all future scheduled appointments. The MIN function finds the nearest date for each patient, and HAVING ensures only patients with confirmed appointments appear in the results.

Query 17:

Counts(*) counts every row in the grouped result, giving the total_emergency_visits for each patient ,then we check if it is ≥ 2 to select patients with at least two emergency visits.MAX(C.Date)finds the most recent date of all emergency visits,which is then compared using DATEDIFF to ensure it is within 14 days of today(using CURDATE())

Query 18:

We need to join the Hospital, Department, Clinical Activity, and Appointment tables to retrieve all the necessary data. Then, we count the number of appointments and rank them by city (so that the results are grouped by city) and by appointment count in descending order.

Query 19:

This query finds medications in each city where the highest and lowest hospital unit prices differ by more than 30%, showing: the city, medication name, minimum/maximum price, and percentage spread between prices. Here's how it works, it joins Hospital, Stock, and Medication to relate city, hospital stock, and medications. For each city and medication, it calculates min/max prices using aggregation. It then Computes price spread percentage: $(\text{max-min})/\text{min} \times 100$. After that, it filters results with HAVING to include only those where the spread is above 30%.

Query 20:

For the data quality check on stock entries, all stock rows with negative quantities or non-positive unit prices were selected. This allowed for detecting potential errors in the stock database that could impact inventory management or reporting

4 Implementation & Results

Validating BCNF, Lossless joins, Dependency preservation:

Normalization

- Validate each relation against BCNF.

1. Patient

FDs:

- $\text{IID} \rightarrow \text{Name, Sex, Birth, BloodGroup, Phone, CIN}$
- $\text{CIN} \rightarrow \text{IID, Name, Sex, Birth, BloodGroup, Phone}$ (CIN is a candidate key)

Both IID and CIN are candidate keys. (superkeys)

Conclusion: BCNF is satisfied

2. ContactLocation

FD: $\text{CLID} \rightarrow \text{Street, City, Province, PostalCode, Phone}$

- CLID is the primary key(supkey)

Conclusion: BCNF is satisfied

3. Staff (with subtypes ISA)

FDs:

- STAFF ID \rightarrow Name, Status
- STAFF ID \rightarrow LicenseNumber, Specialty (Practitioner)
- STAFF ID \rightarrow Grade, Ward (Caregiving)
- STAFF ID \rightarrow Modality, Certifications (Technical)

STAFF ID is the key in all cases.

All Attributes are functionally dependent on STAFF ID.

Conclusion: BCNF is satisfied

4. Hospital

FD: assume HID \rightarrow Name, City, Region

All Attributes are functionally dependent on HID

Conclusion: BCNF is satisfied

5. Department

FDs:

- DEPT ID \rightarrow Name, Specialty, HID

Conclusion: BCNF is satisfied

6. Insurance

FD: InsID \rightarrow Type

Conclusion: BCNF is satisfied

7. ClinicalActivity

FD:

- CAID → Title, Time, Date, IID, STAFF ID, DEPT ID, ExpID

Conclusion: BCNF is satisfied.

8. Expense

FDs:

- ExpID → Total, InsID, CAID
- CAID ↔ ExpID
- Primary key: ExpID
- Transitive: ExpID → CAID → (Patient, Staff, Department, Hospital)

Conclusion: BCNF is satisfied.

9. Prescription

FDs:

- PID → DateIssued

Conclusion: BCNF is satisfied.

10. Appointment (ISA of ClinicalActivity)

FD: CAID → Reason, Status

Conclusion: BCNF is satisfied.

11. Emergency (ISA of ClinicalActivity)

FD: CAID → TriageLevel, Outcome

Conclusion: BCNF is satisfied.

12. Medication

FD: DrugID → Name, Form, Strength, Class, ActiveIngredient, Manufacturer

Conclusion: BCNF is satisfied.

13. Stock (Hospital–Medication)

FD: {HID, DrugID, StockTimestamp} → UnitPrice, StockRemaining, Qty, ReorderLevel

- Key: composite {HID, DrugID, StockTimestamp}

Conclusion: BCNF is satisfied.

14. Prescription–Medication (M:N)

FD: {PID, DrugID} → Dosage, Duration

- Composite key {PID, DrugID}

Conclusion: BCNF is satisfied.

15. Patient–ContactLocation

FD: {IID, CLID} → relationship attributes

- Composite key {IID, CLID}

Conclusion: BCNF is satisfied.

16. Patient–Insurance

FDs: {IID,InsID} → relationship attributes

Conclusion: BCNF is satisfied.

17. Staff–Department (WorkIn)

FD: {STAFF ID,DEPT ID} → (relationship attributes)

Conclusion: BCNF is satisfied.

- Verify lossless joins and dependency preservation after decomposition.

Patient

FDs:

IID → Name,Sex,Birth,BloodGroup,Phone,CIN

CIN → IID

Decomposition:

 $R_1(\text{IID}, \text{Name}, \text{Sex}, \text{Birth}, \text{BloodGroup}, \text{Phone})$ and $R_2(\text{IID}, \text{CIN})$

- Intersection: IID

- IID → R1

Conclusion: Lossless join is satisfied.

For R1:

FD1:

IID → Name,Sex,Birth,BloodGroup,Phone

For R1:

FD1:

CIN → Name,Sex,Birth,BloodGroup,Phone

(FD1 Union FD2)⁺ = FDs⁺

Conclusion: Dependency preserving is satisfied.

NOTE : for the other relations there is no decomposition needed because already their Functional dependencies contain only primarykey→All attributes

We provided alongside this pdf a sql file in the Github Repository named '**Lab5.sql**' that is used to create the database schema.

Inserting 5 samples into each table and then manipulating the data:

```

1 USE lab3;
2
3 INSERT INTO Patient VALUES
4 (1, 'CIN1001', 'Sara\u00d7El\u00d7Amrani', '1995-06-21', 'F', 'A+', '0612345678'),
5 (2, 'CIN1002', 'Omar\u00d7Bennani', '1988-09-15', 'M', 'O-', '0623456789'),

```

```
6 (3, 'CIN1003', 'Youssef\u00a0Khadir', '2000-01-10', 'M', 'B+', '0634567890'),  
7 (4, 'CIN1004', 'Lina\u00a0Berrada', '1993-11-05', 'F', 'AB-', '0645678901'),  
8 (5, 'CIN1005', 'Amine\u00a0Fassi', '1998-04-20', 'M', 'A-', '0656789012');  
9  
10 INSERT INTO Hospital VALUES  
11 (1, 'CHU\u00a0Rabat', 'Rabat', 'Rabat-Sale'),  
12 (2, 'CHU\u00a0Casablanca', 'Casablanca', 'Casablanca-Settat'),  
13 (3, 'CHU\u00a0Marrakech', 'Marrakech', 'Marrakech-Safi'),  
14 (4, 'CHU\u00a0Fes', 'Fes', 'Fes-Meknes'),  
15 (5, 'CHU\u00a0Agadir', 'Agadir', 'Souss-Massa');  
16  
17 INSERT INTO Department VALUES  
18 (1, 1, 'Cardiology', 'Heart\u00a0Diseases'),  
19 (2, 1, 'Neurology', 'Brain\u00a0and\u00a0Nerves'),  
20 (3, 2, 'Pediatrics', 'Child\u00a0Health'),  
21 (4, 3, 'Surgery', 'General\u00a0Surgery'),  
22 (5, 4, 'Radiology', 'Imaging');  
23  
24 INSERT INTO Staff VALUES  
25 (1, 'Dr.\u00a0Hicham\u00a0El\u00a0Idrissi', 'Active'),  
26 (2, 'Dr.\u00a0Salma\u00a0Tazi', 'Active'),  
27 (3, 'Nurse\u00a0Rania\u00a0Lahlou', 'Active'),  
28 (4, 'Technician\u00a0Yassine\u00a0Kabbaj', 'Retired'),  
29 (5, 'Dr.\u00a0Mehdi\u00a0Rahali', 'Active');  
30  
31 INSERT INTO Work_in VALUES  
32 (1, 1),  
33 (2, 2),  
34 (3, 3),  
35 (4, 4),  
36 (4, 5),  
37 (5, 5);  
38  
39 INSERT INTO ClinicalActivity VALUES  
40 (1, 1, 1, 1, '2025-10-10', '10:00:00'),  
41 (2, 2, 2, 2, '2025-10-30', '11:45:00'),  
42 (3, 3, 3, 3, '2025-09-05', '15:00:00'),  
43 (4, 4, 4, 4, '2025-09-22', '12:15:00'),  
44 (5, 5, 4, 5, '2026-01-30', '11:45:00'),  
45 (6, 1, 1, 1, '2025-11-11', '11:45:00'),  
46 (7, 2, 2, 2, '2025-01-11', '09:30:00'),  
47 (8, 5, 5, 5, '2025-11-20', '09:30:00');  
48  
49 INSERT INTO Appointment VALUES  
50 (1, 'Routine\u00a0Checkup', 'Completed'),  
51 (2, 'Neurology\u00a0Consultation', 'Completed'),  
52 (3, 'Pediatric\u00a0Visit', 'Cancelled'),  
53 (4, 'Surgery\u00a0Follow-up', 'Completed'),  
54 (5, 'X-ray\u00a0Review', 'Scheduled'),  
55 (6, 'Routine\u00a0Checkup', 'Completed'),  
56 (7, 'Neurology\u00a0Consultation', 'Completed'),
```

```
57 (8, 'Monitoring\u2014treatment\u2014progress', 'Scheduled');

58
59 INSERT INTO Emergency VALUES
60 (1, 3, 'Discharged'),
61 (2, 4, 'Admitted'),
62 (3, 2, 'Transferred'),
63 (4, 5, 'Deceased'),
64 (5, 1, 'Admitted');

65
66 INSERT INTO Insurance VALUES
67 (1, 'CNOPS'),
68 (2, 'CNSS'),
69 (3, 'RAMED'),
70 (4, 'Private'),
71 (5, 'None');

72
73 INSERT INTO Expense VALUES
74 (1, 1, 1, 500.00),
75 (2, 2, 2, 800.00),
76 (3, 3, 3, 200.00),
77 (4, 4, 4, 1200.00),
78 (5, 5, 5, 50.00);

79
80 INSERT INTO Medication VALUES
81 (1, 'Paracetamol', 'Tablet', '500mg', 'Paracetamol', 'Analgesic', 'Sanofi'),
82 (2, 'Amoxicillin', 'Capsule', '250mg', 'Amoxicillin', 'Antibiotic', 'Pfizer'),
83 (3, 'Aspirin', 'Tablet', '100mg', 'Acetylsalicylic Acid', 'Analgesic', 'Bayer'),
84 (4, 'Cough\u2014Syrup', 'Liquid', '5mg/10ml', 'Dextromethorphan', 'Antitussive',
85 'GSK'),
86 (5, 'Insulin', 'Injection', '10IU/ml', 'Insulin', 'Antidiabetic', 'Novo\u2014
Nordisk');

87
88 INSERT INTO Stock VALUES
89 (1, 1, NOW(), 10.00, 100, 10),
90 (2, 2, NOW(), 25.00, 50, 5),
91 (2, 3, NOW(), 7.00, 200, 20),
92 (3, 4, NOW(), 10.00, 200, 20),
93 (3, 3, NOW(), 5.00, 200, 20),
94 (4, 4, NOW(), 15.00, 30, 10),
95 (5, 5, NOW(), 50.00, 10, 5);

96
97 INSERT INTO Prescription VALUES
98 (1, 1, '2025-10-10'),
99 (2, 2, '2026-01-11'),
100 (3, 3, '2025-09-05'),
101 (4, 4, '2025-09-22'),
102 (5, 5, '2026-01-30');

103
104 INSERT INTO Includes VALUES
105 (1, 1, '1\u2014tablet\u2014every\u20148h', '5\u2014days'),
(2, 2, '1\u2014capsule\u2014every\u201412h', '7\u2014days'),
```

```

106 (3, 3, '1\u_tablet\u_daily', '10\u_days'),
107 (4, 4, '10ml\u_twice\u_daily', '5\u_days'),
108 (5, 5, 'Inject\u_before\u_meals', '30\u_days');

109
110 INSERT INTO ContactLocation VALUES
111 (1, 'Rabat', 'Rabat-Sale', 'Avenue\u_Hassan\u_II', '12', '10000', '0537001122'),
112 (2, 'Casablanca', 'Casablanca-Settat', 'Bd\u_Zerkouni', '22', '20000',
113     '0522003344'),
114 (3, 'Marrakech', 'Marrakech-Safi', 'Rue\u_Mohammed\u_V', '5', '40000',
115     '0524005566'),
116 (4, 'Fes', 'Fes-Meknes', 'Rue\u_Hassan', '8', '30000', '0535007788'),
117 (5, 'Agadir', 'Souss-Massa', 'Bd\u_20\uAout', '3', '80000', '0528009900');

118 INSERT INTO have VALUES
119 (1, 1),
120 (2, 2),
121 (3, 3),
122 (4, 4),
123 (5, 5);

```

Now to alter the tables:

```

1 ALTER TABLE Patient ADD Email VARCHAR(100);
2 UPDATE Patient SET Phone = '0611223344' WHERE IID = 3;
3 UPDATE Hospital SET Region = 'Casablanca-Settat' WHERE HID = 2;
4
5 DELETE FROM Appointment WHERE Status = 'Cancelled';

```

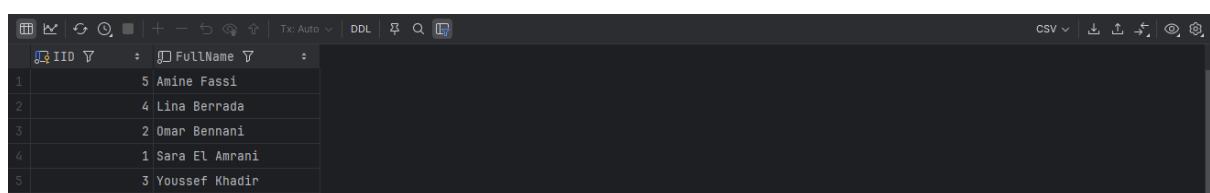
We will now provide screenshots of the results given after running the queries on the database schema created using the "Lab5.sql" file provided alongside this pdf in the GitHub repository after inserting the samples above and manipulation the data.

Query 1:

```

1 SELECT City, COUNT(*) AS total
2 FROM Hospital
3 WHERE City = 'Casablanca'
4 GROUP BY City;

```



IID	FullName
5	Amine Fassi
4	Lina Berrada
2	Omar Bennani
1	Sara El Amrani
3	Youssef Khadir

Query 2:

```

1 SELECT DISTINCT Ins.Type
2 FROM Insurance Ins;

```

Type
1 CNOPS
2 CNSS
3 RAMED
4 Private
5 None

Query 3:

```

1 SELECT DISTINCT S.STAFF_ID , S.FullName
2 FROM Staff S
3 JOIN Work_in W ON S.STAFF_ID = W.STAFF_ID
4 JOIN Department D ON W.DEP_ID = D.DEP_ID
5 JOIN Hospital H ON D.HID = H.HID
6 WHERE H.City='Rabat';

```

STAFF_ID	FullName
1	Dr. Hicham El Idrissi
2	Dr. Salma Tazi

Query 4:

```

1 SELECT A.CAID
2   FROM Appointment A
3 JOIN ClinicalActivity CA ON CA.CAID = A.CAID
4 WHERE CURDATE() <= CA.Date AND DATEDIFF(CA.Date, CURDATE()) <= 7 AND A.Status =
      'Scheduled';

```

A.CAID	Reason	Status	CA.CAID	IID	STAFF_ID	DEP_ID	Date	Time
1	Monitoring treatment progress	Scheduled	8	5	5	5	2025-11-20	09:30:00

Query 5:

```

1 SELECT CA.Dep_ID, COUNT(*) AS cnt
2   FROM Appointment A
3 JOIN ClinicalActivity CA on A.CAID=CA.CAID

```

4 | GROUP BY CA.Dep_ID

Output Result 8 ×

CSV | ↴ ↵ ⌂ ⌃ ⌁ ⌂ ⌃ ⌁

Dep_ID	cnt
1	2
2	2
3	1
4	2

Query 6:

```
1 SELECT S.HID, AVG(UnitPrice) as average_unit_price  
2 FROM Stock S  
3 GROUP BY S.HID;
```

Query 7:

```
1 SELECT HID
2 FROM Hospital H
3 JOIN Department D ON H.HID = D.HID
4 JOIN ClinicalActivity CA ON D.DEP_ID = CA.DEP_ID
5 JOIN Emergency E ON CA.CAID = E.CAID
6 WHERE E.Outcome = 'Admitted'
7 GROUP BY H.HID
8 HAVING COUNT(E.CAID) > 20;
```

The screenshot shows the top ribbon bar of the Microsoft Data Explorer application. It includes standard file and document-related icons (New, Open, Save, Print, etc.), a 'Tx: Auto' dropdown, a 'DDL' button, a search icon, and a refresh icon. On the right side, there are buttons for 'CSV' (with a dropdown arrow), 'Download' (down arrow), 'Upload' (up arrow), 'Copy' (copy icon), 'Paste' (paste icon), and 'Find' (magnifying glass icon). Below the ribbon, there's a toolbar with icons for 'HDFS', 'Hadoop', and other connectivity options.

The result gave an empty table, since we only implemented around 5 samples per table, so no hospital exceed 20 emergency admissions.

Query 8:

```

1 SELECT M.MID, M.Name FROM Medication M
2   JOIN Stock S ON M.MID = S.MID
3 WHERE M.TherapeuticClass = 'Antibiotic' AND S.UnitPrice < 200;

```

MID	Name
1	Amoxicillin

Query 9:

```

1 SELECT S1.HID,S1.Mid
2   FROM Stock S1
3     WHERE (SELECT count(*)
4       FROM Stock S2
5         WHERE S1.HID = S2.HID
6           AND S2.UnitPrice>S1.UnitPrice)<3
7 ORDER BY S1.HID ,S1.UnitPrice DESC;

```

HID	Mid
1	1
2	2
3	2
4	3
5	3
6	4
7	5

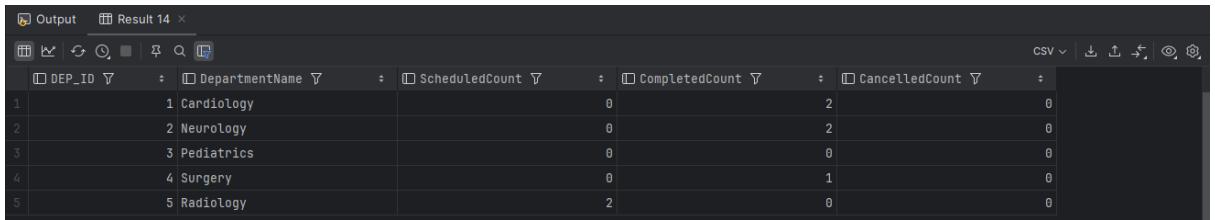
Query 10:

```

1 SELECT D.DEP_ID,
2       D.Name AS DepartmentName,
3       (SELECT COUNT(*)
4         FROM ClinicalActivity CA1
5       JOIN Appointment A1 ON CA1.CAID = A1.CAID
6         WHERE CA1.DEP_ID = D.DEP_ID
7           AND A1.Status = 'Scheduled') AS ScheduledCount,
8       (SELECT COUNT(*)
9         FROM ClinicalActivity CA2
10      JOIN Appointment A2 ON CA2.CAID = A2.CAID
11        WHERE CA2.DEP_ID = D.DEP_ID
12          AND A2.Status = 'Completed') AS CompletedCount,
13       (SELECT COUNT(*)
14         FROM ClinicalActivity CA3
15       JOIN Appointment A3 ON CA3.CAID = A3.CAID
16         WHERE CA3.DEP_ID = D.DEP_ID
17           AND A3.Status = 'Cancelled') AS CancelledCount
18   FROM Department D;

```

Output Result 14 ×



DEP_ID	DepartmentName	ScheduledCount	CompletedCount	CancelledCount
1	Cardiology	0	2	0
2	Neurology	0	2	0
3	Pediatrics	0	0	0
4	Surgery	0	1	0
5	Radiology	2	0	0

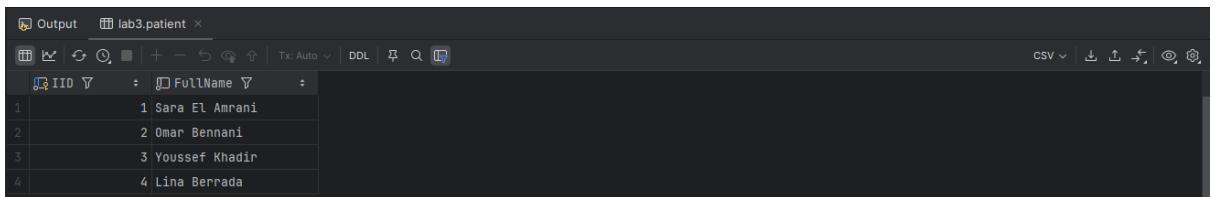
Query 11:

```

1 SELECT P.IID, P.FullName
2   FROM Patient P
3 WHERE P.IID NOT IN (
4     SELECT CA.IID
5       FROM Appointment A
6      JOIN ClinicalActivity CA ON A.CAID = CA.CAID
7     WHERE A.Status = 'Scheduled'
8     AND CA.Date BETWEEN CURDATE() AND DATE_ADD(CURDATE(), INTERVAL 30 DAY)
9 );

```

Output lab3.patient ×



IID	FullName
1	Sara El Amrani
2	Omar Bennani
3	Youssef Khadir
4	Lina Berrada

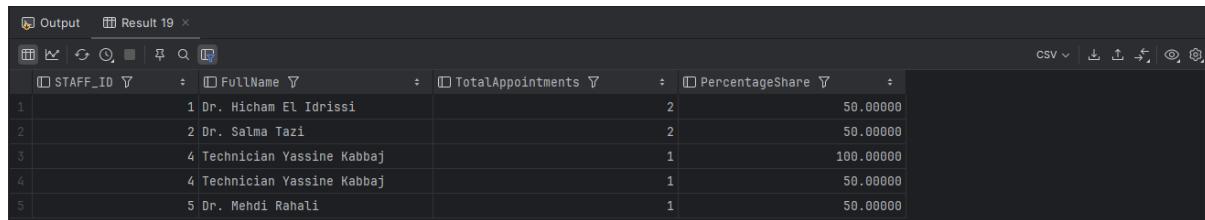
Query 12:

```

1 SELECT S.STAFF_ID,
2       S.FullName,
3       COUNT(A.CAID) AS TotalAppointments,
4       COUNT(A.CAID) * 100.0 /
5         (SELECT COUNT(A2.CAID)
6           FROM ClinicalActivity CA2
7          JOIN Appointment A2 ON CA2.CAID = A2.CAID
8          JOIN Department D2 ON CA2.DEP_ID = D2.DEP_ID
9          WHERE D2.HID = D.HID) AS PercentageShare
10      FROM Staff S
11     JOIN ClinicalActivity CA ON S.STAFF_ID = CA.STAFF_ID
12     JOIN Appointment A ON CA.CAID = A.CAID
13     JOIN Department D ON D.DEP_ID = CA.DEP_ID
14 GROUP BY S.STAFF_ID, S.FullName;

```

Output Result 19 ×



A screenshot of a database query result titled "Result 19". The table has three columns: STAFF_ID, FullName, TotalAppointments, and PercentageShare. The data is as follows:

STAFF_ID	FullName	TotalAppointments	PercentageShare
1	Dr. Wiam El Idrissi	2	50.00000
2	Dr. Salma Tazi	2	50.00000
3	Technician Yassine Kabbaj	1	100.00000
4	Technician Yassine Kabbaj	1	50.00000
5	Dr. Mehdi Rahali	1	50.00000

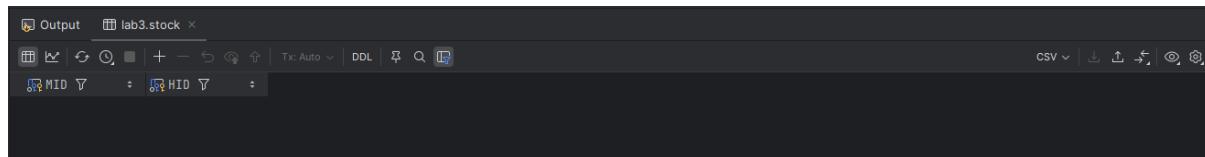
Query 13:

```

1 SELECT MID, HID
2   FROM Stock S
3 WHERE S.ReorderLevel > S.Qty
4 ORDER BY MID;

```

Output lab3.stock ×



A screenshot of a database query result titled "lab3.stock". The table has two columns: MID and HID. The data is as follows:

MID	HID

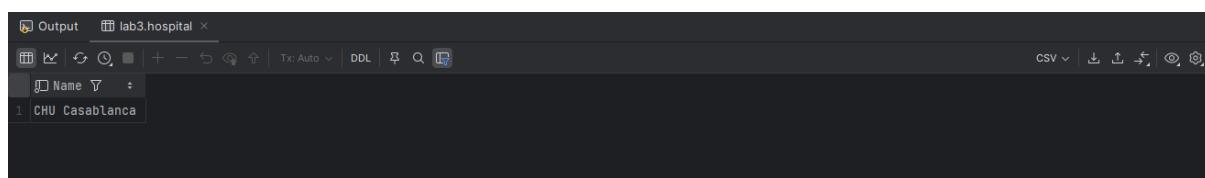
Query 14:

```

1 SELECT H.Name
2   FROM Hospital H
3 JOIN Stock S ON S.HID = H.HID
4 JOIN Medication M ON S.MID = M.MID
5 WHERE NOT EXISTS ((SELECT MID
6                   FROM Medication
7                  WHERE TherapeuticClass = 'Antibiotic')
8 EXCEPT
9           (SELECT S2.MID
10              FROM Stock S2
11             JOIN Medication M2 ON M2.MID = S2.MID
12            WHERE S2.HID = H.HID AND M2.TherapeuticClass =
13              'Antibiotic'));

```

Output lab3.hospital ×



A screenshot of a database query result titled "lab3.hospital". The table has one column: Name. The data is as follows:

Name
CHU Casablanca

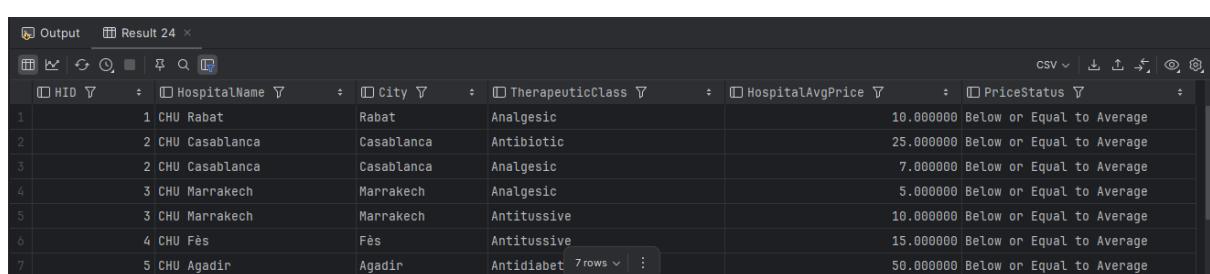
Query 15:

```

1 SELECT H.HID, H.Name AS HospitalName , H.City , M.TherapeuticClass ,
2     AVG(S.UnitPrice) AS HospitalAvgPrice , 'Above Average' AS PriceStatus
3 FROM Hospital H, Stock S, Medication M
4 WHERE H.HID = S.HID
5 AND S.MID = M.MID
6 GROUP BY H.HID, M.TherapeuticClass
7 HAVING AVG(S.UnitPrice) > (
8     SELECT AVG(S2.UnitPrice)
9     FROM Stock S2, Hospital H2, Medication M2
10    WHERE S2.HID = H2.HID
11    AND S2.MID = M2.MID
12    AND H2.City = H.City
13    AND M2.TherapeuticClass = M.TherapeuticClass
14 )
15 UNION
16
17 SELECT H.HID, H.Name AS HospitalName , H.City , M.TherapeuticClass ,
18     AVG(S.UnitPrice) AS HospitalAvgPrice , 'Below or Equal to Average' AS
19     PriceStatus
20 FROM Hospital H, Stock S, Medication M
21 WHERE H.HID = S.HID
22 AND S.MID = M.MID
23 GROUP BY H.HID, M.TherapeuticClass
24 HAVING AVG(S.UnitPrice) <= (
25     SELECT AVG(S2.UnitPrice)
26     FROM Stock S2, Hospital H2, Medication M2
27    WHERE S2.HID = H2.HID
28    AND S2.MID = M2.MID
29    AND H2.City = H.City
30    AND M2.TherapeuticClass = M.TherapeuticClass
31 );

```

Output Result 24 ×



	HID	HospitalName	City	TherapeuticClass	HospitalAvgPrice	PriceStatus
1	1	CHU Rabat	Rabat	Analgesic	10.000000	Below or Equal to Average
2	2	CHU Casablanca	Casablanca	Antibiotic	25.000000	Below or Equal to Average
3	2	CHU Casablanca	Casablanca	Analgesic	7.000000	Below or Equal to Average
4	3	CHU Marrakech	Marrakech	Analgesic	5.000000	Below or Equal to Average
5	3	CHU Marrakech	Marrakech	Antitussive	10.000000	Below or Equal to Average
6	4	CHU Fès	Fès	Antitussive	15.000000	Below or Equal to Average
7	5	CHU Agadir	Agadir	Antidiabet	50.000000	Below or Equal to Average

Query 16:

```

1 SELECT P.IID, P.FullName, MIN(CA.Date) AS NextAppointmentDate
2 FROM Patient P
3 JOIN ClinicalActivity CA ON P.IID = CA.IID
4 JOIN Appointment A ON CA.CAID = A.CAID

```

```

5 WHERE CA.Date >= CURDATE()
6 AND A.Status = 'Scheduled'
7 GROUP BY P.IID
8 HAVING MIN(CA.Date) IS NOT NULL;

```

	IID	FullName	NextAppointmentDate
1	5	Amine Fassi	2025-11-20

Query 17:

```

1 SELECT P.IID,P.FullName,MAX(C.Date)as latest_emergency_date,COUNT(*) as
2   total_emergency_visit
3 FROM Patient P
4 JOIN ClinicalActivity C on P.IID=C.IID
5 JOIN Emergency E on C.CAID=E.CAID
6 GROUP BY P.IID
Having count(*)>=2 AND DATEDIFF(CURDATE(), MAX(C.Date)) <= 14;

```

	IID	FullName	latest_emergency_date	total_emergency_visit

As expected, no patient showed up, and this matches the sample data used in the database.

Query 18:

```

1 SELECT H.City, H.Name, COUNT(DISTINCT A.CAID) as app_count
2 FROM Hospital H
3 JOIN Department D ON H.HID = D.HID
4 JOIN ClinicalActivity CA ON CA.DEP_ID = D.DEP_ID
5 JOIN Appointment A ON A.CAID = CA.CAID
6 WHERE CA.Date BETWEEN CURDATE() AND DATE_ADD(CURDATE(), INTERVAL 90 DAY)
7 AND A.Status = 'Completed'
8 GROUP BY H.HID
9 ORDER BY H.City, app_count;

```

	City	Name	app_count
1	Marrakech	CHU Marrakech	1
2	Rabat	CHU Rabat	3

Query 19:

```

1  SELECT
2      H.City,
3      M.MID,
4      M.Name AS MedicationName,
5      MIN(ST.UnitPrice) AS MinPrice,
6      MAX(ST.UnitPrice) AS MaxPrice,
7      ROUND((MAX(ST.UnitPrice) - MIN(ST.UnitPrice)) * 100.0 / MIN(ST.UnitPrice),
8            2) AS PriceSpreadPercent
9  FROM Hospital H
10 JOIN Stock ST ON H.HID = ST.HID
11 JOIN Medication M ON ST.MID = M.MID
12 GROUP BY H.City, M.MID, M.Name
13 HAVING (MAX(ST.UnitPrice) - MIN(ST.UnitPrice)) * 100.0 / MIN(ST.UnitPrice)>30;

```

In this example, the query doesn't return much information since only around five sample entries were added to each table.

Not a lot of samples were inserted onto the tables, thus, the min and max range are always the same as seen in the image below when replacing $> 30\%$ with $\geq 0\%$

1	City	⋮	MID	⋮	MedicationName	⋮	MinPrice	⋮	MaxPrice	⋮	PriceSpreadPercent
1	Robat				1 Paracetamol		10.00		10.00		0.00
2	Casablanca				2 Amoxicillin		25.00		25.00		0.00
3	Casablanca				3 Aspirin		7.00		7.00		0.00
4	Marrakech				3 Aspirin		5.00		5.00		0.00
5	Marrakech				4 Cough Syrup		10.00		10.00		0.00
6	Fès				4 Cough Syrup		15.00		15.00		0.00
7	Agadir				5 Insulin			7 rows			50.00
								⋮	0		0.00

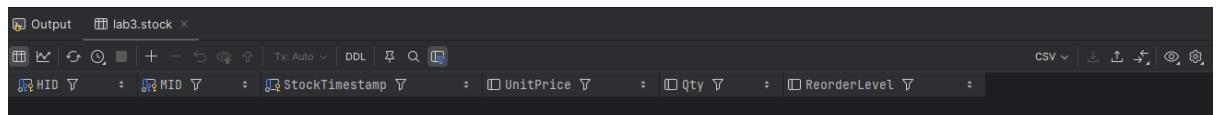
```

1  SELECT
2      H.City,
3      M.MID,
4      M.Name AS MedicationName,
5      MIN(ST.UnitPrice) AS MinPrice,
6      MAX(ST.UnitPrice) AS MaxPrice,
7      ROUND((MAX(ST.UnitPrice) - MIN(ST.UnitPrice)) * 100.0 / MIN(ST.UnitPrice),
8            2) AS PriceSpreadPercent
9  FROM Hospital H
10 JOIN Stock ST ON H.HID = ST.HID
11 JOIN Medication M ON ST.MID = M.MID
12 GROUP BY H.City, M.MID, M.Name
13 HAVING (MAX(ST.UnitPrice) - MIN(ST.UnitPrice)) * 100.0 / MIN(ST.UnitPrice)>=0;

```

Query 20:

```
1 SELECT *
2   FROM Stock
3 WHERE Qty < 0 OR UnitPrice <= 0;
```



Since all inserted values were within the expected range and none were negative, the query produces no output.

5 Discussion

Challenges Faced

2 During the implementation of these SQL queries, we encountered several technical challenges that required careful problem-solving. One significant difficulty was managing complex multi-table joins across four or more tables, as seen in Queries 7 and 18, where connecting Hospital through Department to ClinicalActivity and then to Emergency or Appointment tables demanded precise join conditions to avoid incorrect results. Another challenge emerged in queries requiring division operations, such as Query 14, where we had to identify hospitals stocking all antibiotics using NOT EXISTS patterns.

Another area that required particular attention was working with correlated subqueries for analytical calculations. For instance, in Query 12, calculating the percentage share of appointments per hospital while correctly grouping results by staff members proved challenging. We had to carefully structure the query to reference outer query elements like D.HID within subqueries without creating circular dependencies. Furthermore, queries involving date-based filtering, such as Query 16 for upcoming appointments and Query 17 for recent emergency visits, demanded precise use of date functions and interval calculations to ensure accurate results for operational planning.

Observations

Throughout the query development process, we observed several important patterns in healthcare data analysis. Queries involving patient scheduling and appointments (Queries 4, 10, 11, 16) consistently required joining the ClinicalActivity and Appointment tables with careful status filtering. We noticed that medication-related queries (Queries 8, 13, 15, 19) formed a distinct category, typically joining Medication with Stock and Hospital tables to analyze pricing, availability, and distribution patterns.

The aggregation patterns varied significantly based on the analysis objective. Queries like 5 and 7 used straightforward COUNT aggregations, while more complex analyses in Queries 12 and 15 required nested aggregations with subqueries. We also observed that queries dealing with absence or exclusion criteria, such as Query 11 finding patients without appointments, effectively used NOT IN and subquery patterns rather than outer joins with NULL checks.

Lessons Learned

This project taught us useful skills for working with healthcare data. We learned to split complex questions into simpler SQL queries, which made them easier to write and understand. For example, Query 15 showed us how to use UNION to organize hospitals into different price categories, while other queries helped us compare medication costs across different locations.

We also discovered that some query structures work better than others when dealing with large amounts of data. Most importantly, we found that well-planned queries don't just answer one question - they can be reused for future analysis. The work we did here creates a strong base for the hospital system to make better day-to-day decisions and plan for the future.

6 Conclusion

In this part of the project, we applied the principles of database normalization and SQL query formulation to the MNHS dataset. Normalisation allowed us to refine the relational schema by removing redundancy, enforcing data consistency, and ensuring that all relations respected BCNF properties. This guarantees a clean and reliable database foundation for analytical operations. Using the refined schema, we implemented a completed set of SQL queries covering simple selections, multi-table relationships, nested subqueries, aggregate functions, and conditional counting. These queries demonstrate how normalized data can be effectively exploited to extract meaningful insights. Overall, this lab highlights the importance of combining a well-designed schema with expressive SQL queries to support healthcare data management.